

Introduction

Customer churn, the phenomenon where customers stop using a company's products or services, is a significant challenge for many businesses. Understanding and predicting churn can help companies take proactive steps to retain customers, thereby reducing loss of revenue and improving long-term profitability. This project aims to predict customer churn by analyzing behavioral and demographic data using machine learning techniques. By leveraging a comprehensive dataset, the project identifies key factors contributing to churn and develops models to accurately predict which customers are at risk of leaving.

Data Source

The primary dataset was sourced from a telecommunications company, encompassing various customer attributes and their churn status. The dataset includes 7,043 rows and 21 columns, with a binary target variable indicating whether the customer has churned.

Data Source: <https://www.kaggle.com/datasets/blastchar/telco-customer-churn>

Necessary Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objects as go
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier
from sklearn import metrics
from sklearn.metrics import roc_curve
from sklearn.metrics import recall_score, confusion_matrix, precision_score, f1_score, accuracy_score, classification_report
```

```
In [3]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.regularizers import l2
```

Data Exploration

```
In [4]: df = pd.read_csv("Telco-Customer-Churn.csv")
```

```
In [5]: df.head()
```

Out[5]:	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	De
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 21 columns

In [6]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
In [7]: df.describe()
```

Out[7]:

	SeniorCitizen	tenure	MonthlyCharges
--	---------------	--------	----------------

count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

In [8]: `df.isnull().sum()`

Out[8]:

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	0
Churn	0
dtype:	int64

```
In [9]: df.shape
```

```
Out[9]: (7043, 21)
```

```
In [10]: # Iterate over all columns in the DataFrame
for column in df.columns:
    print(f"Value counts for column: {column}")
    print(df[column].value_counts())
    print("\n" + "-"*50 + "\n")
```

```
Value counts for column: customerID
customerID
7590-VHVEG      1
3791-LGQCY      1
6008-NAIXK      1
5956-YHHRX      1
5365-LLFYV      1
..
9796-MVYXX      1
2637-FKFSY      1
1552-AAGRXX     1
4304-TSPVK      1
3186-AJIEK      1
Name: count, Length: 7043, dtype: int64
```

```
-----

Value counts for column: gender
gender
Male      3555
Female    3488
Name: count, dtype: int64
```

```
-----

Value counts for column: SeniorCitizen
SeniorCitizen
0      5901
1      1142
Name: count, dtype: int64
```

Value counts for column: Partner

Partner

No 3641

Yes 3402

Name: count, dtype: int64

Value counts for column: Dependents

Dependents

No 4933

Yes 2110

Name: count, dtype: int64

Value counts for column: tenure

tenure

1 613

72 362

2 238

3 200

4 176

...

28 57

39 56

44 51

36 50

0 11

Name: count, Length: 73, dtype: int64

Value counts for column: PhoneService

PhoneService

Yes 6361

No 682

Name: count, dtype: int64

Value counts for column: MultipleLines

MultipleLines

No	3390
----	------

Yes	2971
-----	------

No phone service	682
------------------	-----

Name: count, dtype: int64

Value counts for column: InternetService

InternetService

Fiber optic	3096
-------------	------

DSL	2421
-----	------

No	1526
----	------

Name: count, dtype: int64

Value counts for column: OnlineSecurity

OnlineSecurity

No	3498
----	------

Yes	2019
-----	------

No internet service	1526
---------------------	------

Name: count, dtype: int64

Value counts for column: OnlineBackup

OnlineBackup

No	3088
----	------

Yes	2429
-----	------

No internet service	1526
---------------------	------

Name: count, dtype: int64

Value counts for column: DeviceProtection

DeviceProtection


```
No                3095
Yes               2422
No internet service  1526
Name: count, dtype: int64
```

Value counts for column: TechSupport

TechSupport

```
No                3473
Yes               2044
No internet service  1526
Name: count, dtype: int64
```

Value counts for column: StreamingTV

StreamingTV

```
No                2810
Yes               2707
No internet service  1526
Name: count, dtype: int64
```

Value counts for column: StreamingMovies

StreamingMovies

```
No                2785
Yes               2732
No internet service  1526
Name: count, dtype: int64
```

Value counts for column: Contract

Contract

```
Month-to-month    3875
Two year          1695
One year          1473
Name: count, dtype: int64
```

Value counts for column: PaperlessBilling
PaperlessBilling
Yes 4171
No 2872
Name: count, dtype: int64

Value counts for column: PaymentMethod
PaymentMethod
Electronic check 2365
Mailed check 1612
Bank transfer (automatic) 1544
Credit card (automatic) 1522
Name: count, dtype: int64

Value counts for column: MonthlyCharges
MonthlyCharges
20.05 61
19.85 45
19.95 44
19.90 44
20.00 43
 ..
23.65 1
114.70 1
43.65 1
87.80 1
78.70 1
Name: count, Length: 1585, dtype: int64

Value counts for column: TotalCharges
TotalCharges
 11
20.2 11

```
19.75      9
20.05      8
19.9       8
..
6849.4     1
692.35     1
130.15     1
3211.9     1
6844.5     1
Name: count, Length: 6531, dtype: int64
```


```
Value counts for column: Churn
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

Data Visualization


```
In [11]: # Data preparation
 churn_counts = df['Churn'].value_counts()
 labels = churn_counts.index
 sizes = churn_counts.values
 colors = ['#ff9999', '#66b3ff']

 plt.figure(figsize=(8, 8))
 plt.pie(sizes, labels=labels, colors=colors, autopct='%1.1f%%', startangle=140, wedgeprops={'edgecolor': 'black'})
 plt.title('Percentage of Churn', fontsize=16)
 plt.show()
```

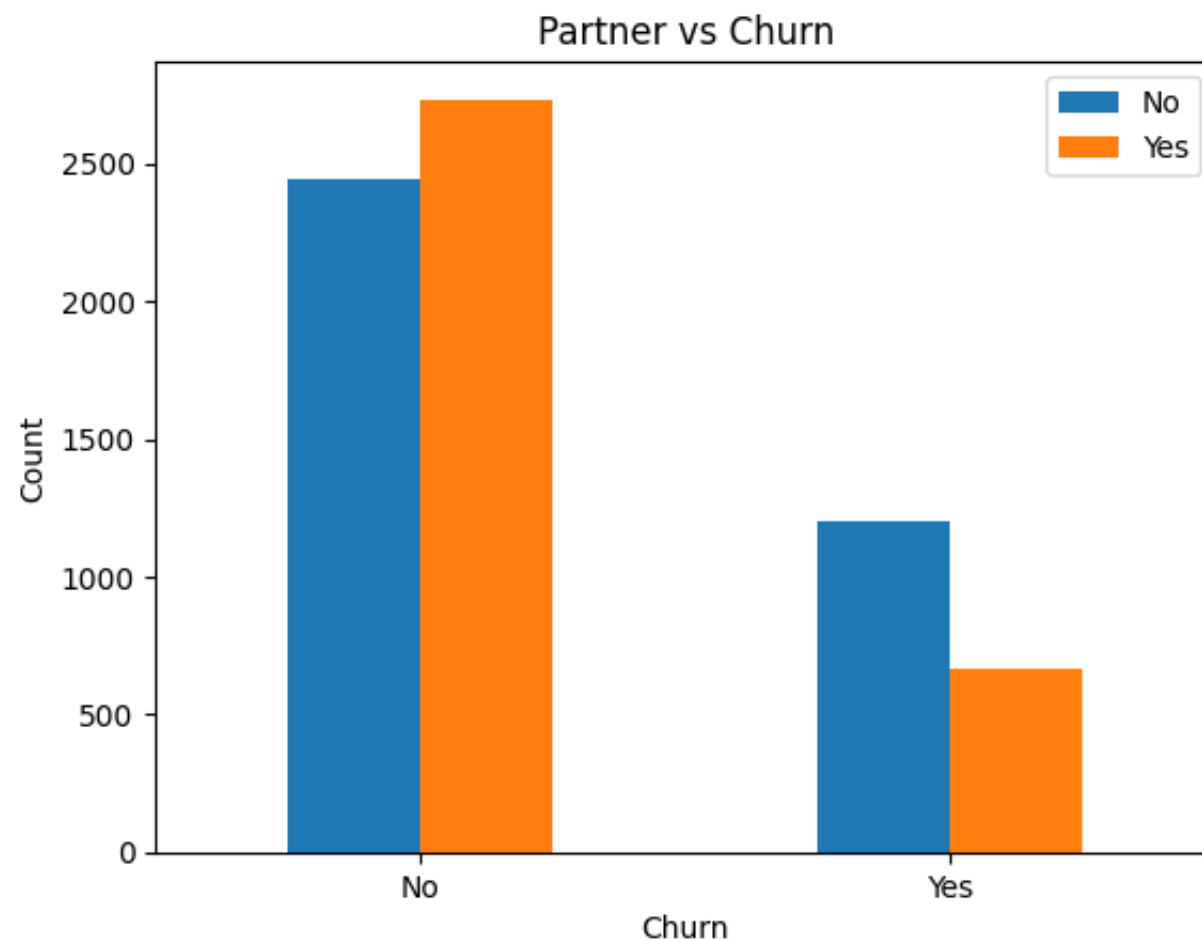
Percentage of Churn



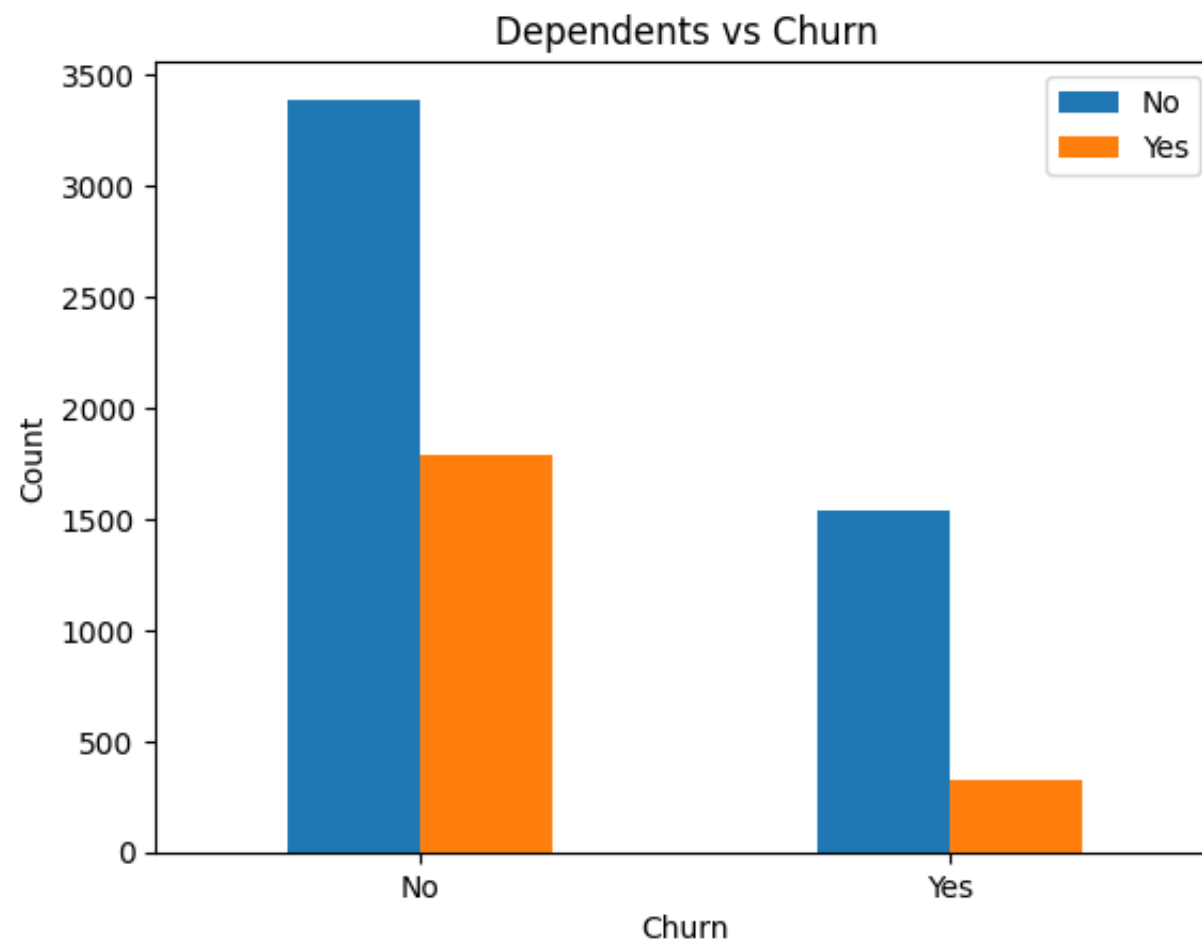
```
In [12]: pd.crosstab(df['Churn'], df['gender']).plot(kind='bar')
plt.title('Gender vs Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.legend(['Female', 'Male'])
plt.xticks(rotation=0)
plt.show()
```



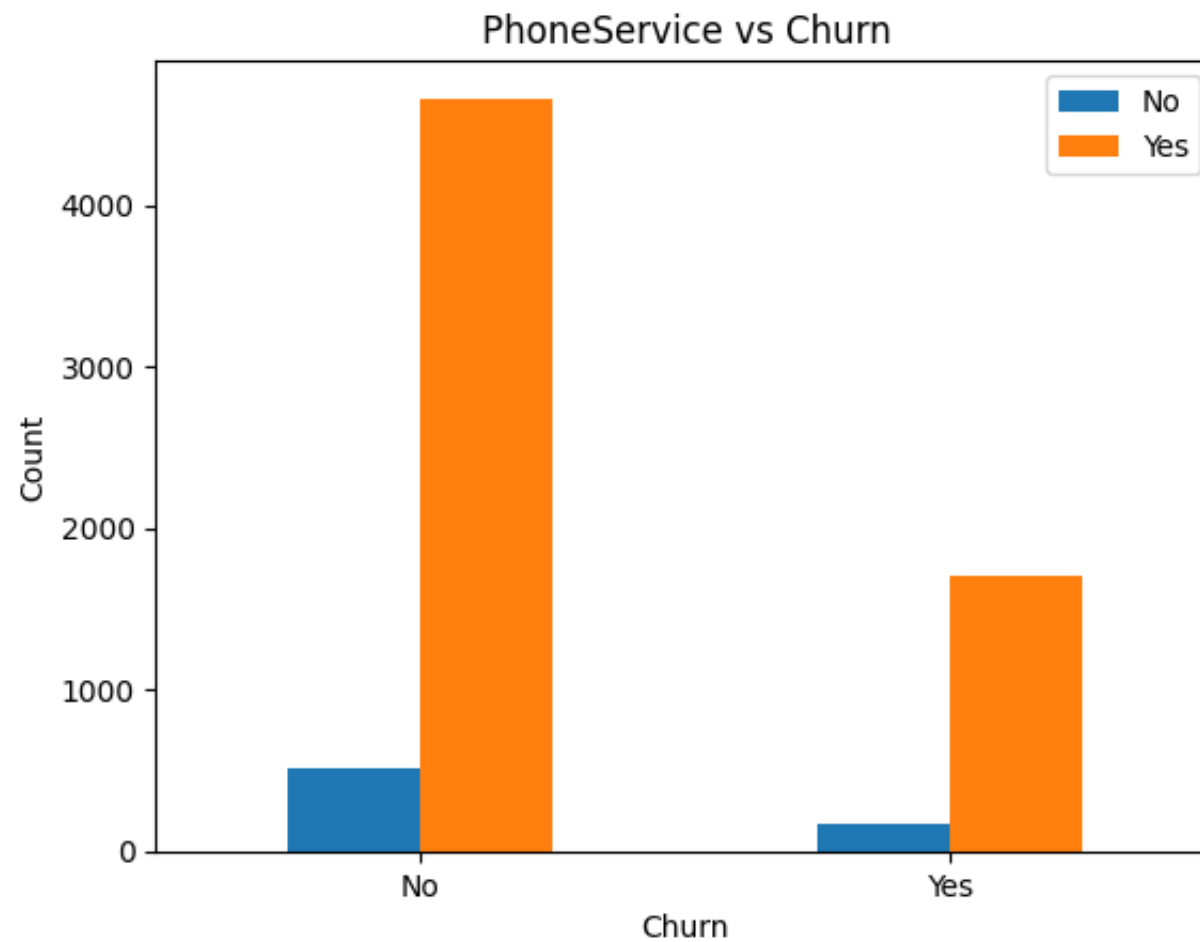
```
In [13]: pd.crosstab(df['Churn'], df['Partner']).plot(kind='bar')
plt.title('Partner vs Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.legend(['No', 'Yes'])
plt.xticks(rotation=0)
plt.show()
```



```
In [14]: pd.crosstab(df['Churn'], df['Dependents']).plot(kind='bar')
plt.title('Dependents vs Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.legend(['No', 'Yes'])
plt.xticks(rotation=0)
plt.show()
```



```
In [15]: pd.crosstab(df['Churn'], df['PhoneService']).plot(kind='bar')
plt.title('PhoneService vs Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.legend(['No', 'Yes'])
plt.xticks(rotation=0)
plt.show()
```




```
In [16]: # Create a crosstab of Churn and MultipleLines
churn_multiplelines = pd.crosstab(df['Churn'], df['MultipleLines'])

# Reset the index for easier plotting
churn_multiplelines = churn_multiplelines.reset_index()

# Melt the DataFrame for Seaborn compatibility
churn_multiplelines_melted = churn_multiplelines.melt(id_vars='Churn', value_vars=churn_multiplelines.columns[1:])

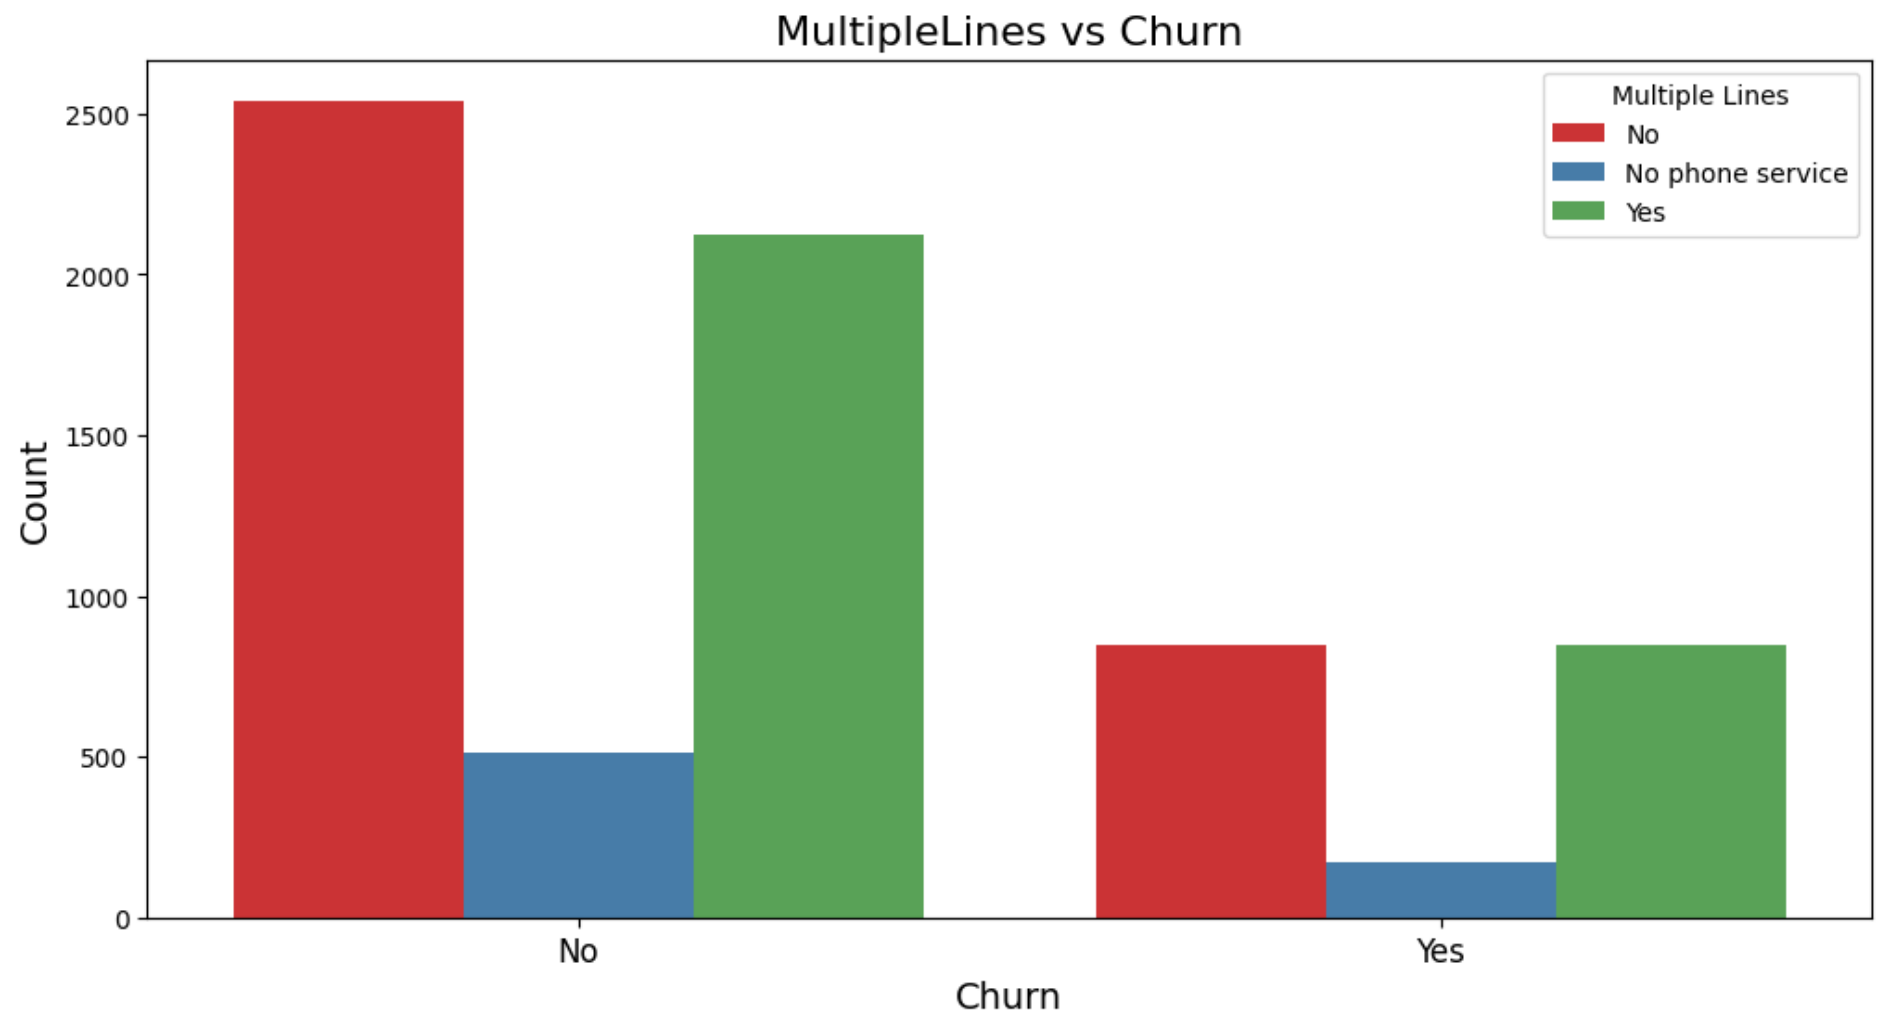
# Plot using Seaborn
plt.figure(figsize=(12, 6))
sns.barplot(x='Churn', y='Count', hue='MultipleLines', data=churn_multiplelines_melted, palette='Set1')

# Add title and labels
plt.title('MultipleLines vs Churn', fontsize=16)
plt.xlabel('Churn', fontsize=14)
plt.ylabel('Count', fontsize=14)

# Customize legend
plt.legend(title='Multiple Lines')


# Adjust x-ticks for better readability
plt.xticks(rotation=0, fontsize=12)

# Display the plot
plt.show()
```




```
In [17]: fig = px.histogram(df, x="Churn", color="Contract", barmode="group", title="<b>Customer contract distribution<b>"  
fig.show()
```

Customer contract distribution



```
In [18]: fig = px.histogram(df, x="Churn", color="PaymentMethod", barmode="group", title="<b>Customer payment method distribution</b>")
fig.show()
```

Customer payment method distribution



```
In [19]: import plotly.graph_objects as go
fig = go.Figure()

fig.add_trace(go.Bar(
    x = ['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [965, 992, 219, 240],
    name = 'DSL',
))


fig.add_trace(go.Bar(
    x = ['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [889, 910, 664, 633],
    name = 'Fiber optic',
))

fig.add_trace(go.Bar(
    x = ['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
        ["Female", "Male", "Female", "Male"]],
    y = [690, 717, 56, 57],
    name = 'No Internet',
))

fig.update_layout(title_text="<b>Churn Distribution w.r.t. Internet Service and Gender</b>")


fig.show()
```

Churn Distribution w.r.t. Internet Service and Gender




```
In [20]: fig = px.histogram(df, x="Churn", color="InternetService", barmode="group", title="<b>Customer internet service c
fig.show()
```

Customer internet service distribution

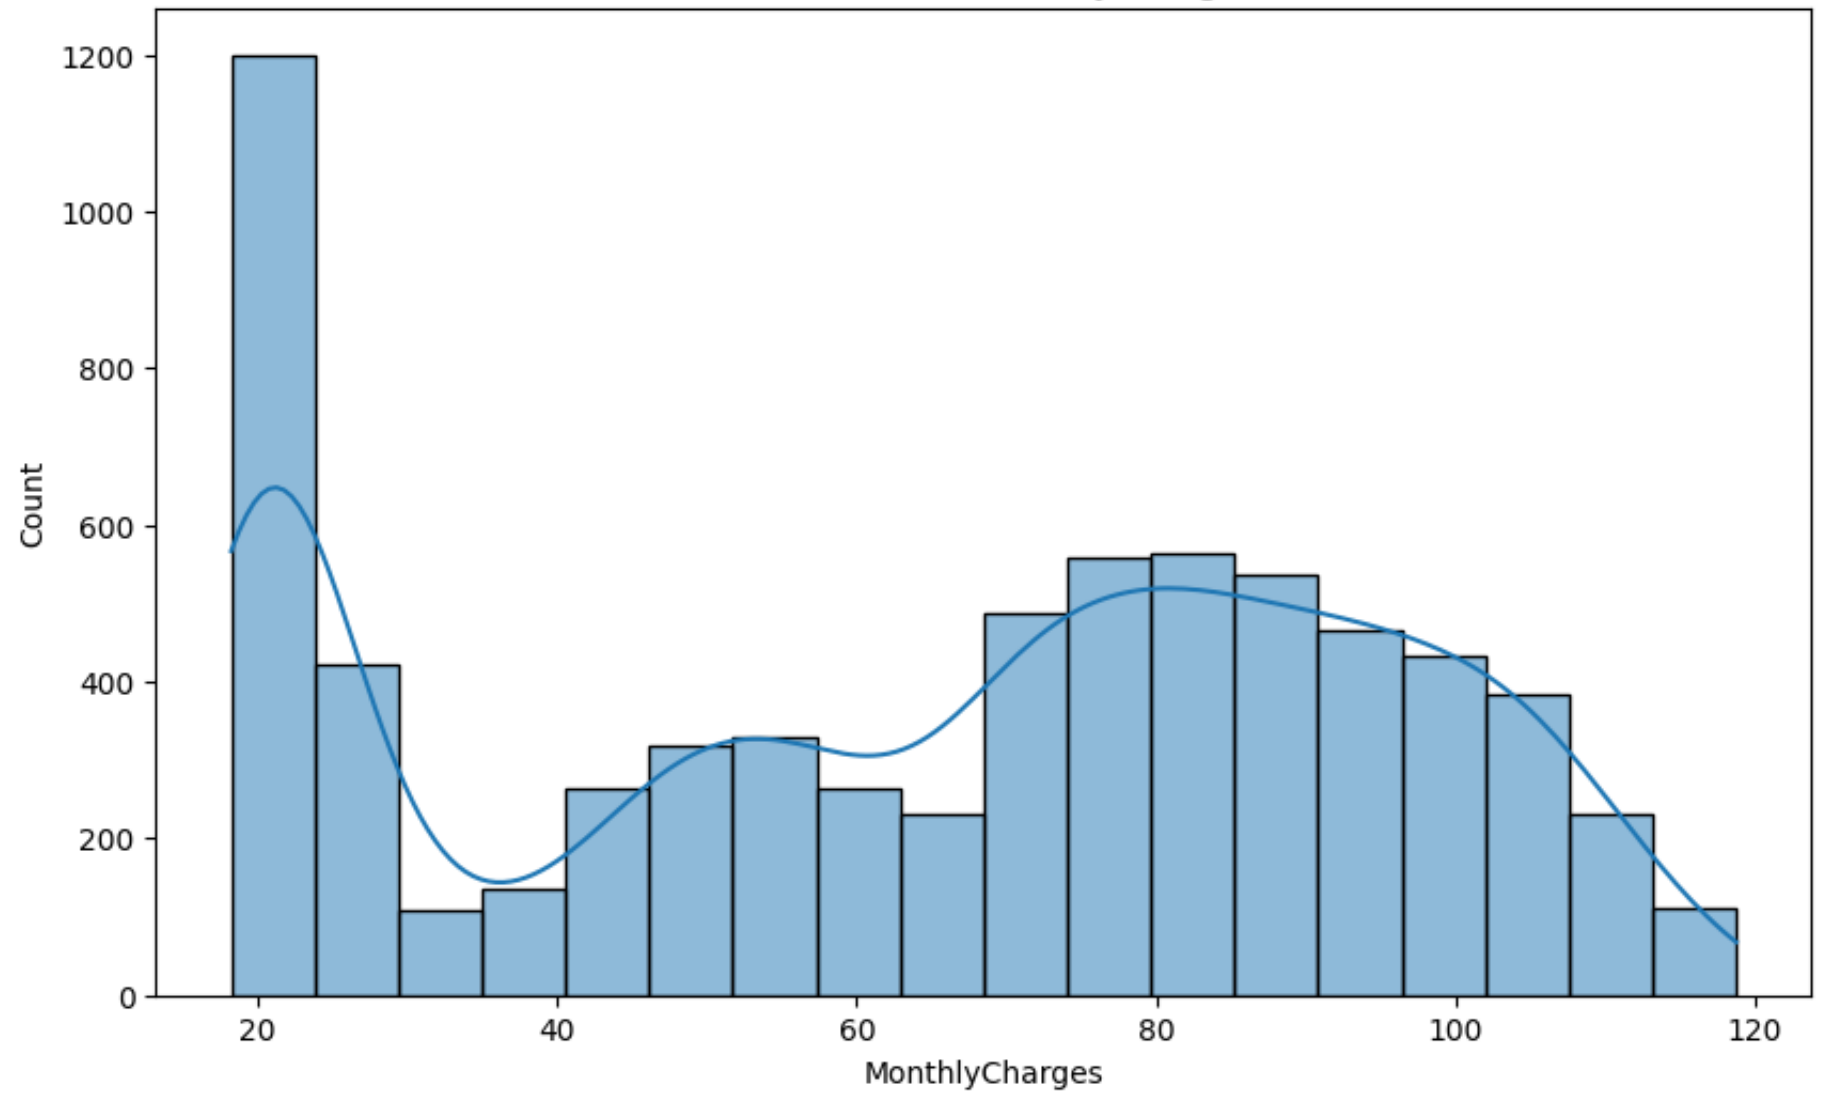


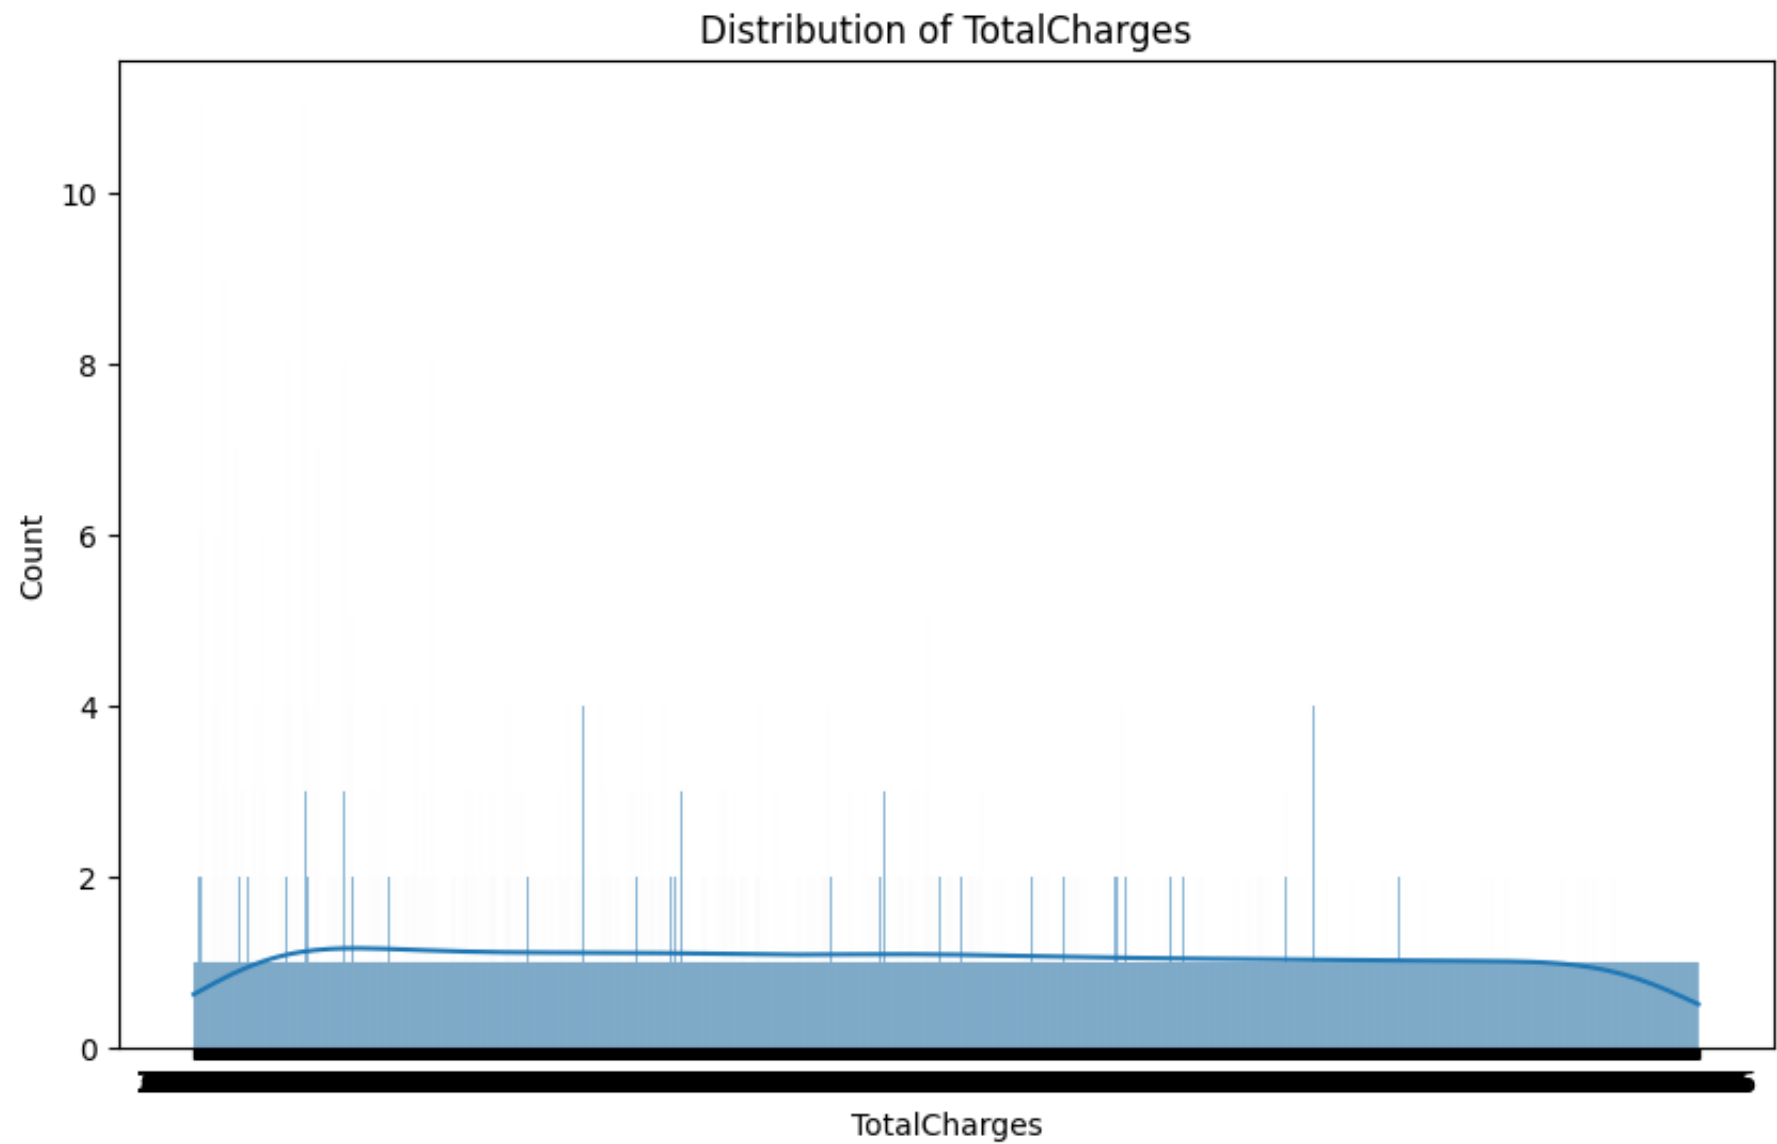
```
In [21]: num_cols = ["tenure", "MonthlyCharges", "TotalCharges"]
for col in num_cols:
    plt.figure(figsize=(10, 6))
    sns.histplot(df[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.show()
```

Distribution of tenure



Distribution of MonthlyCharges





Data Handling

```
In [22]: df.drop('customerID', axis=1, inplace=True)  
df.head()
```

Out[22]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Devi
0	Female	0	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	0	No	No	34	Yes	No	DSL	Yes	No	
2	Male	0	No	No	2	Yes	No	DSL	Yes	Yes	
3	Male	0	No	No	45	No	No phone service	DSL	Yes	No	
4	Female	0	No	No	2	Yes	No	Fiber optic	No	No	

In [23]:

```
df['TotalCharges'] = pd.to_numeric(df.TotalCharges, errors='coerce')
df.isnull().sum()
```

Out[23]:

gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0
dtype:	int64

```
In [24]: df.fillna(df["TotalCharges"].mean(),inplace=True)
df.isnull().sum()
```

```
Out[24]: gender                0
SeniorCitizen                0
Partner                      0
Dependents                   0
tenure                       0
PhoneService                 0
MultipleLines                0
InternetService              0
OnlineSecurity               0
OnlineBackup                 0
DeviceProtection             0
TechSupport                  0
StreamingTV                  0
StreamingMovies              0
Contract                     0
PaperlessBilling             0
PaymentMethod                0
MonthlyCharges               0
TotalCharges                  0
Churn                         0
dtype: int64
```

```
In [25]: df["SeniorCitizen"] = df["SeniorCitizen"].map({0: "No", 1: "Yes"})
df.head()
```

Out [25]:	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Devi
0	Female	No	Yes	No	1	No	No phone service	DSL	No	Yes	
1	Male	No	No	No	34	Yes	No	DSL	Yes	No	
2	Male	No	No	No	2	Yes	No	DSL	Yes	Yes	
3	Male	No	No	No	45	No	No phone service	DSL	Yes	No	
4	Female	No	No	No	2	Yes	No	Fiber optic	No	No	

In [26]: `df["InternetService"].describe(include=['object', 'bool'])`

Out [26]:

```

count          7043
unique           3
top      Fiber optic
freq          3096
Name: InternetService, dtype: object

```

In [27]: `df.describe(exclude='object')`

Out [27]:

	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	2283.300441
std	24.559481	30.090047	2265.000258
min	0.000000	18.250000	18.800000
25%	9.000000	35.500000	402.225000
50%	29.000000	70.350000	1400.550000
75%	55.000000	89.850000	3786.600000
max	72.000000	118.750000	8684.800000

```
In [28]: # Set the context for the plot
sns.set_context("paper", font_scale=1.2)

# Set the figure size
plt.figure(figsize=(10, 6))


# Plot KDE for "Not Churn"
sns.kdeplot(df.MonthlyCharges[df["Churn"] == 'No'], shade=True, color="blue", label="Not Churn", bw_adjust=1.2)

# Plot KDE for "Churn"
sns.kdeplot(df.MonthlyCharges[df["Churn"] == 'Yes'], shade=True, color="red", label="Churn", bw_adjust=1.2)

# Customize the title and labels
plt.title('Distribution of Monthly Charges by Churn', fontsize=16, weight='bold')
plt.xlabel('Monthly Charges', fontsize=14)
plt.ylabel('Density', fontsize=14)

# Customize the legend
plt.legend(title='Churn Status', loc='upper right', fontsize=12, title_fontsize='13')

# Display the plot
plt.show()
```



Data Preprocessing

```
In [29]: def object_to_int(dataframe_series):
         if dataframe_series.dtype=='object':
             dataframe_series = LabelEncoder().fit_transform(dataframe_series)
         return dataframe_series
```

```
In [30]: df = df.apply(lambda x: object_to_int(x))
         df.head()
```

```
Out[30]:
```


	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	Devi
0	0	0	1	0	1	0	1	0	0	2	
1	1	0	0	0	34	1	0	0	2	0	
2	1	0	0	0	2	1	0	0	2	2	
3	1	0	0	0	45	0	1	0	2	0	
4	0	0	0	0	2	1	0	1	0	0	

```
In [31]: X = df.iloc[:, :-1]
         y = df.iloc[:, -1]
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```


```
In [32]: def distplot(feature, frame, color='g'):
         plt.figure(figsize=(12, 7))
         plt.title("Distribution for {}".format(feature))
         ax = sns.distplot(frame[feature], color= color)
```

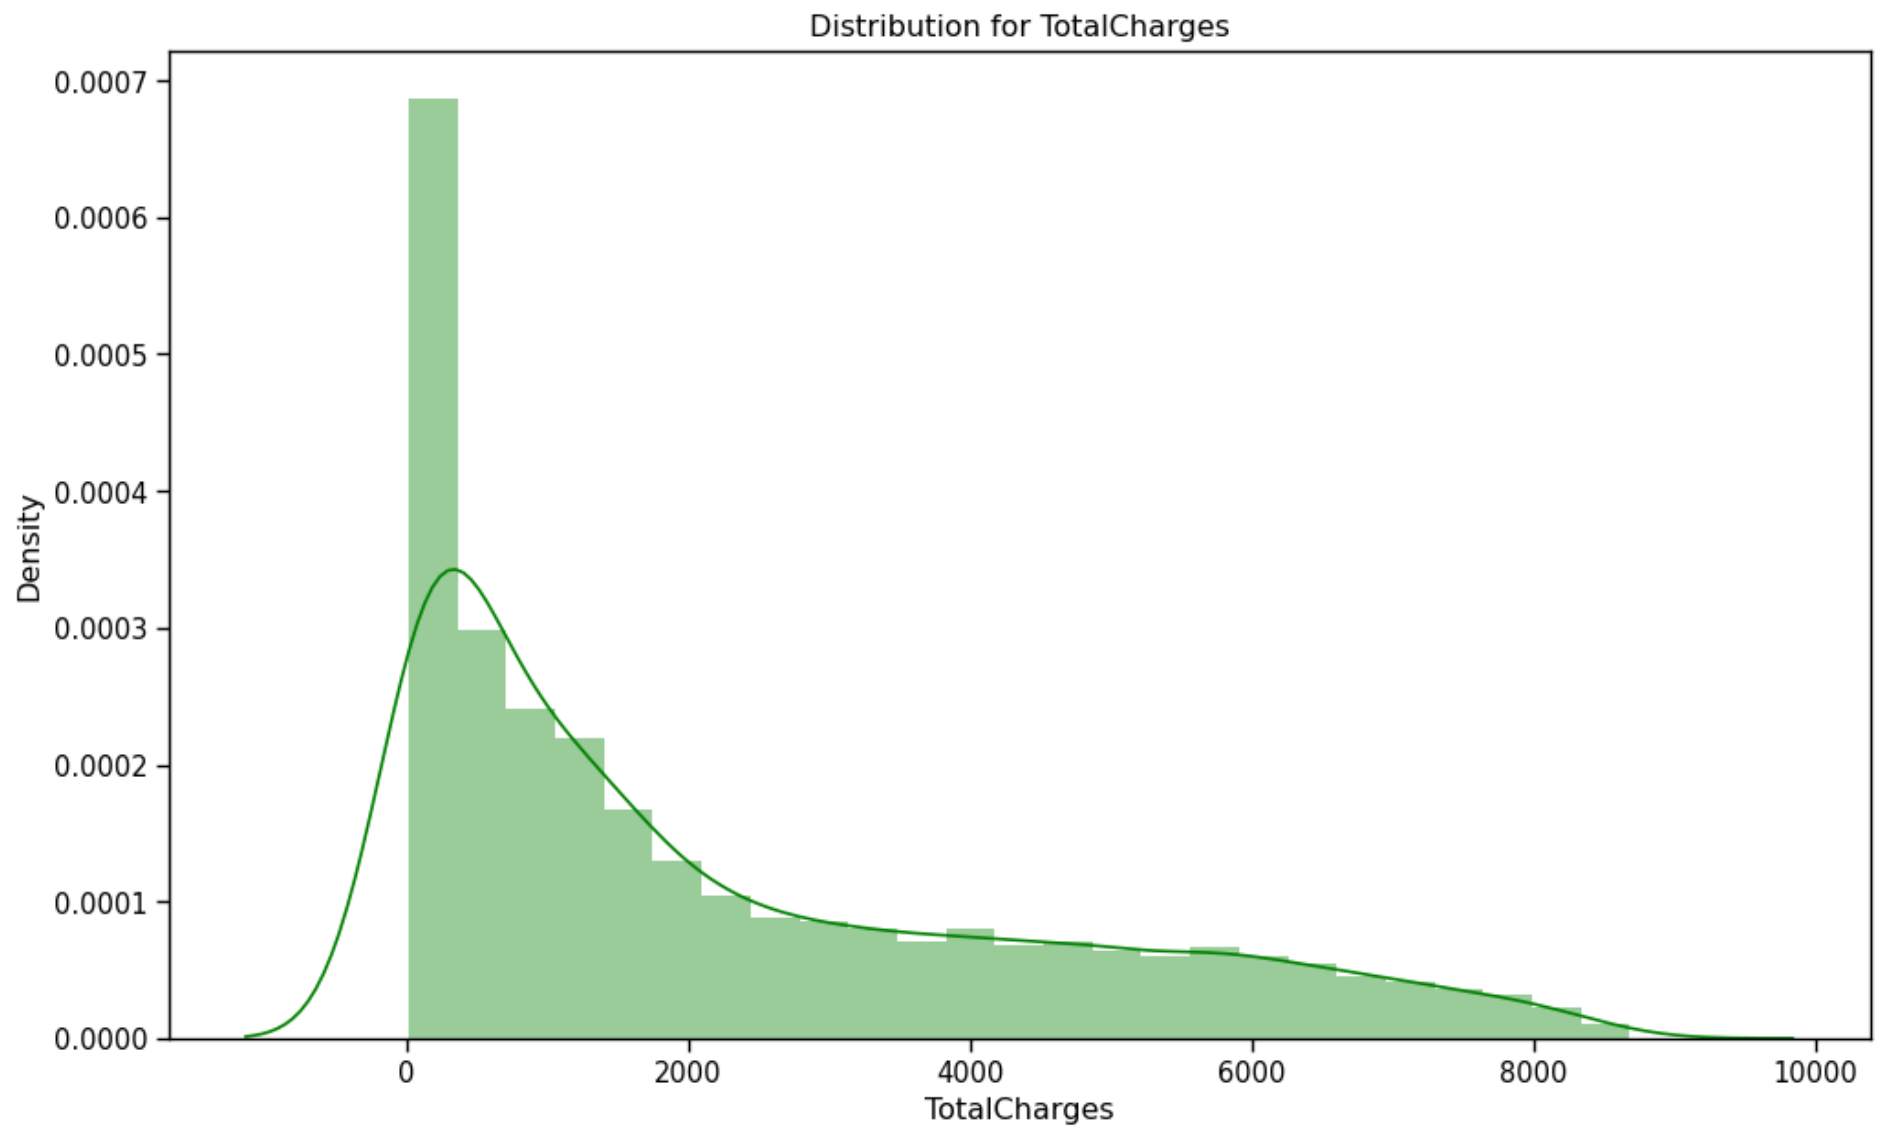
```
In [33]: num_cols = ["tenure", 'MonthlyCharges', 'TotalCharges']
         for feat in num_cols: distplot(feat, df)
```


Distribution for tenure



Distribution for MonthlyCharges





```
In [34]: # StandardScaler
scaler= StandardScaler()

X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

Machine Learning Models

KNN

```
In [35]: knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train, y_train)
predicted_y = knn_model.predict(X_test)
accuracy_knn = knn_model.score(X_test, y_test)
print("KNN accuracy:", accuracy_knn)
```

KNN accuracy: 0.7700496806245565

```
In [36]: print(classification_report(y_test, predicted_y))
```

	precision	recall	f1-score	support
0	0.79	0.94	0.86	1036
1	0.64	0.31	0.41	373
accuracy			0.77	1409
macro avg	0.71	0.62	0.64	1409
weighted avg	0.75	0.77	0.74	1409

Random Forest

```
In [37]: model_rf = RandomForestClassifier()
model_rf.fit(X_train, y_train)


# Make predictions
prediction_test = model_rf.predict(X_test)
print(metrics.accuracy_score(y_test, prediction_test))
```

0.7970191625266146

```
In [38]: print(classification_report(y_test, prediction_test))
```

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.66	0.49	0.56	373
accuracy			0.80	1409
macro avg	0.74	0.70	0.71	1409
weighted avg	0.79	0.80	0.79	1409

```
In [39]: y_rfpred_prob = model_rf.predict_proba(X_test)[:,-1]
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_rfpred_prob)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rf, tpr_rf, label='Random Forest',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve',fontsize=16)
plt.show();
```



Logistic Regression

```
In [40]: lr_model = LogisticRegression()
lr_model.fit(X_train,y_train)
accuracy_lr = lr_model.score(X_test,y_test)
print("Logistic Regression accuracy is :",accuracy_lr)
```

Logistic Regression accuracy is : 0.8168914123491838

```
In [41]: lr_pred= lr_model.predict(X_test)
print(classification_report(y_test,lr_pred))
```

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1036
1	0.68	0.58	0.63	373
accuracy			0.82	1409
macro avg	0.77	0.74	0.75	1409
weighted avg	0.81	0.82	0.81	1409

XGBOOST

```
In [42]: xgb_model = XGBClassifier()
xgb_model.fit(X_train, y_train)
accuracy_xgb = xgb_model.score(X_test,y_test)
print("XGBOOST accuracy is :",accuracy_xgb)
```

XGBOOST accuracy is : 0.7821149751596878

```
In [43]: print(classification_report(y_test, xgb_model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.83	0.88	0.86	1036
1	0.61	0.50	0.55	373
accuracy			0.78	1409
macro avg	0.72	0.69	0.70	1409
weighted avg	0.77	0.78	0.77	1409

```
In [44]: # KNN
```

```
knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train, y_train)
predicted_y = knn_model.predict(X_test)
accuracy_knn = knn_model.score(X_test,y_test)
```

```

print("KNN accuracy:", accuracy_knn)

print(classification_report(y_test, predicted_y))

# Random Forest

model_rf = RandomForestClassifier()
model_rf.fit(X_train, y_train)

# Make predictions
prediction_test = model_rf.predict(X_test)
print(metrics.accuracy_score(y_test, prediction_test))

print(classification_report(y_test, prediction_test))

y_rfpred_prob = model_rf.predict_proba(X_test)[:,1]
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_rfpred_prob)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rf, tpr_rf, label='Random Forest', color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve', fontsize=16)
plt.show();

# Logistic Regression

lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
accuracy_lr = lr_model.score(X_test, y_test)
print("Logistic Regression accuracy is :", accuracy_lr)

lr_pred = lr_model.predict(X_test)
print(classification_report(y_test, lr_pred))

# XGBOOST

xgb_model = XGBClassifier()
xgb_model.fit(X_train, y_train)
accuracy_xgb = xgb_model.score(X_test, y_test)
print("XGBOOST accuracy is :", accuracy_xgb)

```



```
print(classification_report(y_test, xgb_model.predict(X_test)))
```


KNN accuracy: 0.7700496806245565

	precision	recall	f1-score	support
0	0.79	0.94	0.86	1036
1	0.64	0.31	0.41	373
accuracy			0.77	1409
macro avg	0.71	0.62	0.64	1409
weighted avg	0.75	0.77	0.74	1409

0.7955997161107168

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.65	0.48	0.56	373
accuracy			0.80	1409
macro avg	0.74	0.70	0.71	1409
weighted avg	0.78	0.80	0.78	1409

Random Forest ROC Curve



Logistic Regression accuracy is : 0.8168914123491838

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1036
1	0.68	0.58	0.63	373
accuracy			0.82	1409
macro avg	0.77	0.74	0.75	1409
weighted avg	0.81	0.82	0.81	1409

XGBOOST accuracy is : 0.7821149751596878

	precision	recall	f1-score	support
0	0.83	0.88	0.86	1036
1	0.61	0.50	0.55	373
accuracy			0.78	1409
macro avg	0.72	0.69	0.70	1409
weighted avg	0.77	0.78	0.77	1409

Deep Learning Model

In [45]: `model = Sequential()`

```
# Input layer with L2 regularization
model.add(Dense(128, input_dim=X_train.shape[1], activation='relu', kernel_regularizer=l2(0.001)))

# Additional hidden layers with L2 regularization and Dropout
model.add(Dense(64, activation='relu', kernel_regularizer=l2(0.0001)))
model.add(Dropout(0.3))
model.add(Dense(32, activation='relu', kernel_regularizer=l2(0.0001)))
model.add(Dropout(0.3))
model.add(Dense(16, activation='relu', kernel_regularizer=l2(0.0001)))

# Output layer
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.0001), loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(X_train, y_train, epochs=100, batch_size=64, validation_data=(X_test, y_test))
```

Epoch 1/100

89/89 ————— 1s 3ms/step - accuracy: 0.5850 - loss: 0.7115 - val_accuracy: 0.7353 - val_loss: 0.5935

Epoch 2/100

89/89 ————— 0s 1ms/step - accuracy: 0.7241 - loss: 0.6051 - val_accuracy: 0.7353 - val_loss: 0.5296

Epoch 3/100

89/89 ————— 0s 1ms/step - accuracy: 0.7309 - loss: 0.5609 - val_accuracy: 0.7402 - val_loss: 0.4982

Epoch 4/100

89/89 ————— 0s 1ms/step - accuracy: 0.7414 - loss: 0.5287 - val_accuracy: 0.7743 - val_loss: 0.4820

Epoch 5/100

89/89 ————— 0s 1ms/step - accuracy: 0.7611 - loss: 0.5169 - val_accuracy: 0.8006 - val_loss: 0.4729

Epoch 6/100


89/89 ————— 0s 1ms/step - accuracy: 0.7665 - loss: 0.5098 - val_accuracy: 0.8077 - val_loss: 0.4665

Epoch 7/100

89/89 ————— 0s 1ms/step - accuracy: 0.7836 - loss: 0.4964 - val_accuracy: 0.8133 - val_loss: 0.46

40

Epoch 8/100

89/89  **0s** 1ms/step - accuracy: 0.7695 - loss: 0.4986 - val_accuracy: 0.8126 - val_loss: 0.46


07

Epoch 9/100

89/89  **0s** 2ms/step - accuracy: 0.7843 - loss: 0.4990 - val_accuracy: 0.8148 - val_loss: 0.45

81

Epoch 10/100

89/89  **0s** 2ms/step - accuracy: 0.7812 - loss: 0.5031 - val_accuracy: 0.8105 - val_loss: 0.45

65

Epoch 11/100

89/89  **0s** 1ms/step - accuracy: 0.7775 - loss: 0.4906 - val_accuracy: 0.8169 - val_loss: 0.45

48

Epoch 12/100

89/89  **0s** 1ms/step - accuracy: 0.7782 - loss: 0.5024 - val_accuracy: 0.8155 - val_loss: 0.45

48

Epoch 13/100

89/89  **0s** 1ms/step - accuracy: 0.7846 - loss: 0.4954 - val_accuracy: 0.8133 - val_loss: 0.45


29

Epoch 14/100

89/89  **0s** 1ms/step - accuracy: 0.7845 - loss: 0.4858 - val_accuracy: 0.8148 - val_loss: 0.45


17

Epoch 15/100

89/89  **0s** 1ms/step - accuracy: 0.7895 - loss: 0.4824 - val_accuracy: 0.8119 - val_loss: 0.45

10

Epoch 16/100

89/89  **0s** 1ms/step - accuracy: 0.7891 - loss: 0.4737 - val_accuracy: 0.8133 - val_loss: 0.44

91

Epoch 17/100

89/89  **0s** 1ms/step - accuracy: 0.7932 - loss: 0.4808 - val_accuracy: 0.8112 - val_loss: 0.44


91

Epoch 18/100

89/89  **0s** 986us/step - accuracy: 0.7881 - loss: 0.4838 - val_accuracy: 0.8112 - val_loss: 0.

4486

Epoch 19/100

89/89  **0s** 1ms/step - accuracy: 0.7861 - loss: 0.4838 - val_accuracy: 0.8148 - val_loss: 0.44















79


Epoch 20/100


89/89  **0s** 1ms/step - accuracy: 0.8024 - loss: 0.4637 - val_accuracy: 0.8148 - val_loss: 0.44


77


Epoch 21/100


89/89  0s 1ms/step - accuracy: 0.7903 - loss: 0.4877 - val_accuracy: 0.8162 - val_loss: 0.44
60
Epoch 22/100
89/89  0s 1ms/step - accuracy: 0.7888 - loss: 0.4773 - val_accuracy: 0.8141 - val_loss: 0.44
59
Epoch 23/100
89/89  0s 1ms/step - accuracy: 0.7988 - loss: 0.4648 - val_accuracy: 0.8155 - val_loss: 0.44
62
Epoch 24/100
89/89  0s 1ms/step - accuracy: 0.7957 - loss: 0.4668 - val_accuracy: 0.8162 - val_loss: 0.44
45
Epoch 25/100
89/89  0s 1ms/step - accuracy: 0.7876 - loss: 0.4791 - val_accuracy: 0.8141 - val_loss: 0.44
43
Epoch 26/100
89/89  0s 985us/step - accuracy: 0.7970 - loss: 0.4691 - val_accuracy: 0.8176 - val_loss: 0.
4433
Epoch 27/100
89/89  0s 1ms/step - accuracy: 0.8046 - loss: 0.4621 - val_accuracy: 0.8169 - val_loss: 0.44
39
Epoch 28/100
89/89  0s 1ms/step - accuracy: 0.7976 - loss: 0.4777 - val_accuracy: 0.8190 - val_loss: 0.44
24
Epoch 29/100
89/89  0s 1ms/step - accuracy: 0.7972 - loss: 0.4698 - val_accuracy: 0.8183 - val_loss: 0.44
20
Epoch 30/100
89/89  0s 1ms/step - accuracy: 0.7904 - loss: 0.4647 - val_accuracy: 0.8190 - val_loss: 0.44
14
Epoch 31/100
89/89  0s 1ms/step - accuracy: 0.7976 - loss: 0.4623 - val_accuracy: 0.8197 - val_loss: 0.44
23
Epoch 32/100
89/89  0s 1ms/step - accuracy: 0.7986 - loss: 0.4636 - val_accuracy: 0.8204 - val_loss: 0.44
09
Epoch 33/100
89/89  0s 1ms/step - accuracy: 0.7995 - loss: 0.4641 - val_accuracy: 0.8190 - val_loss: 0.43
97
Epoch 34/100
89/89  0s 1ms/step - accuracy: 0.7975 - loss: 0.4625 - val_accuracy: 0.8204 - val_loss: 0.43
96


Epoch 35/100
89/89  0s 1ms/step - accuracy: 0.8106 - loss: 0.4524 - val_accuracy: 0.8190 - val_loss: 0.4402


Epoch 36/100
89/89  0s 1ms/step - accuracy: 0.7991 - loss: 0.4669 - val_accuracy: 0.8190 - val_loss: 0.4392


Epoch 37/100
89/89  0s 1ms/step - accuracy: 0.8037 - loss: 0.4589 - val_accuracy: 0.8211 - val_loss: 0.4394


Epoch 38/100
89/89  0s 1ms/step - accuracy: 0.8024 - loss: 0.4626 - val_accuracy: 0.8197 - val_loss: 0.4398


Epoch 39/100
89/89  0s 1ms/step - accuracy: 0.8060 - loss: 0.4640 - val_accuracy: 0.8169 - val_loss: 0.4384


Epoch 40/100
89/89  0s 1ms/step - accuracy: 0.8002 - loss: 0.4515 - val_accuracy: 0.8204 - val_loss: 0.4389


Epoch 41/100
89/89  0s 1ms/step - accuracy: 0.7950 - loss: 0.4568 - val_accuracy: 0.8197 - val_loss: 0.4383


Epoch 42/100
89/89  0s 1ms/step - accuracy: 0.7998 - loss: 0.4536 - val_accuracy: 0.8197 - val_loss: 0.4378


Epoch 43/100
89/89  0s 1ms/step - accuracy: 0.8025 - loss: 0.4537 - val_accuracy: 0.8183 - val_loss: 0.4390

Epoch 44/100
89/89  0s 1ms/step - accuracy: 0.7985 - loss: 0.4641 - val_accuracy: 0.8190 - val_loss: 0.4377

Epoch 45/100
89/89  0s 1ms/step - accuracy: 0.8007 - loss: 0.4584 - val_accuracy: 0.8219 - val_loss: 0.4378


Epoch 46/100
89/89  0s 1ms/step - accuracy: 0.8042 - loss: 0.4524 - val_accuracy: 0.8233 - val_loss: 0.4378

Epoch 47/100
89/89  0s 1ms/step - accuracy: 0.8024 - loss: 0.4536 - val_accuracy: 0.8219 - val_loss: 0.4379

Epoch 48/100
89/89  0s 1ms/step - accuracy: 0.8025 - loss: 0.4536 - val_accuracy: 0.8219 - val_loss: 0.43

75

Epoch 49/100

89/89  **0s** 1ms/step - accuracy: 0.8031 - loss: 0.4508 - val_accuracy: 0.8197 - val_loss: 0.43

77

Epoch 50/100

89/89  **0s** 1ms/step - accuracy: 0.8066 - loss: 0.4498 - val_accuracy: 0.8211 - val_loss: 0.43


77

Epoch 51/100

89/89  **0s** 1ms/step - accuracy: 0.8047 - loss: 0.4575 - val_accuracy: 0.8226 - val_loss: 0.43


67

Epoch 52/100

89/89  **0s** 1ms/step - accuracy: 0.8010 - loss: 0.4634 - val_accuracy: 0.8219 - val_loss: 0.43

74

Epoch 53/100

89/89  **0s** 994us/step - accuracy: 0.8070 - loss: 0.4485 - val_accuracy: 0.8204 - val_loss: 0.

4359

Epoch 54/100

89/89  **0s** 1ms/step - accuracy: 0.7984 - loss: 0.4560 - val_accuracy: 0.8219 - val_loss: 0.43

65

Epoch 55/100

89/89  **0s** 1ms/step - accuracy: 0.8047 - loss: 0.4507 - val_accuracy: 0.8226 - val_loss: 0.43


72

Epoch 56/100

89/89  **0s** 998us/step - accuracy: 0.8063 - loss: 0.4521 - val_accuracy: 0.8211 - val_loss: 0.


4359

Epoch 57/100

89/89  **0s** 993us/step - accuracy: 0.8044 - loss: 0.4466 - val_accuracy: 0.8211 - val_loss: 0.


4358

Epoch 58/100

89/89  **0s** 995us/step - accuracy: 0.8089 - loss: 0.4552 - val_accuracy: 0.8240 - val_loss: 0.


4352

Epoch 59/100

89/89  **0s** 998us/step - accuracy: 0.7983 - loss: 0.4631 - val_accuracy: 0.8254 - val_loss: 0.


4355

Epoch 60/100

89/89  **0s** 1ms/step - accuracy: 0.8105 - loss: 0.4298 - val_accuracy: 0.8240 - val_loss: 0.43















59


Epoch 61/100


89/89  **0s** 1ms/step - accuracy: 0.8019 - loss: 0.4418 - val_accuracy: 0.8261 - val_loss: 0.43


54


Epoch 62/100


89/89  0s 1ms/step - accuracy: 0.8128 - loss: 0.4406 - val_accuracy: 0.8247 - val_loss: 0.4351
Epoch 63/100
89/89  0s 1ms/step - accuracy: 0.8072 - loss: 0.4427 - val_accuracy: 0.8247 - val_loss: 0.4355
Epoch 64/100
89/89  0s 1ms/step - accuracy: 0.8037 - loss: 0.4447 - val_accuracy: 0.8247 - val_loss: 0.4367
Epoch 65/100
89/89  0s 1ms/step - accuracy: 0.8168 - loss: 0.4406 - val_accuracy: 0.8247 - val_loss: 0.4361
Epoch 66/100
89/89  0s 992us/step - accuracy: 0.8103 - loss: 0.4388 - val_accuracy: 0.8240 - val_loss: 0.4351
Epoch 67/100
89/89  0s 991us/step - accuracy: 0.8050 - loss: 0.4431 - val_accuracy: 0.8240 - val_loss: 0.4341
Epoch 68/100
89/89  0s 1ms/step - accuracy: 0.8085 - loss: 0.4469 - val_accuracy: 0.8204 - val_loss: 0.4333
Epoch 69/100
89/89  0s 1ms/step - accuracy: 0.8109 - loss: 0.4320 - val_accuracy: 0.8226 - val_loss: 0.4336
Epoch 70/100
89/89  0s 1ms/step - accuracy: 0.8124 - loss: 0.4383 - val_accuracy: 0.8219 - val_loss: 0.4333
Epoch 71/100
89/89  0s 1ms/step - accuracy: 0.8043 - loss: 0.4430 - val_accuracy: 0.8233 - val_loss: 0.4329
Epoch 72/100
89/89  0s 1ms/step - accuracy: 0.8174 - loss: 0.4318 - val_accuracy: 0.8233 - val_loss: 0.4327
Epoch 73/100
89/89  0s 1ms/step - accuracy: 0.8142 - loss: 0.4384 - val_accuracy: 0.8226 - val_loss: 0.4328
Epoch 74/100
89/89  0s 1ms/step - accuracy: 0.8141 - loss: 0.4405 - val_accuracy: 0.8211 - val_loss: 0.4330
Epoch 75/100
89/89  0s 1ms/step - accuracy: 0.8239 - loss: 0.4285 - val_accuracy: 0.8240 - val_loss: 0.4336


Epoch 76/100
89/89  0s 1ms/step - accuracy: 0.8140 - loss: 0.4365 - val_accuracy: 0.8219 - val_loss: 0.4329


Epoch 77/100
89/89  0s 1ms/step - accuracy: 0.8028 - loss: 0.4402 - val_accuracy: 0.8233 - val_loss: 0.4330


Epoch 78/100
89/89  0s 1ms/step - accuracy: 0.8103 - loss: 0.4445 - val_accuracy: 0.8211 - val_loss: 0.4326


Epoch 79/100
89/89  0s 1ms/step - accuracy: 0.8178 - loss: 0.4313 - val_accuracy: 0.8211 - val_loss: 0.4334


Epoch 80/100
89/89  0s 1ms/step - accuracy: 0.8109 - loss: 0.4406 - val_accuracy: 0.8226 - val_loss: 0.4339


Epoch 81/100
89/89  0s 1ms/step - accuracy: 0.8093 - loss: 0.4413 - val_accuracy: 0.8211 - val_loss: 0.4331


Epoch 82/100
89/89  0s 1ms/step - accuracy: 0.8144 - loss: 0.4321 - val_accuracy: 0.8197 - val_loss: 0.4333


Epoch 83/100
89/89  0s 1ms/step - accuracy: 0.8180 - loss: 0.4370 - val_accuracy: 0.8197 - val_loss: 0.4336


Epoch 84/100
89/89  0s 1ms/step - accuracy: 0.8110 - loss: 0.4452 - val_accuracy: 0.8211 - val_loss: 0.4334

Epoch 85/100
89/89  0s 1ms/step - accuracy: 0.8209 - loss: 0.4312 - val_accuracy: 0.8133 - val_loss: 0.4341

Epoch 86/100
89/89  0s 1ms/step - accuracy: 0.8153 - loss: 0.4333 - val_accuracy: 0.8204 - val_loss: 0.4335

Epoch 87/100
89/89  0s 1ms/step - accuracy: 0.8108 - loss: 0.4419 - val_accuracy: 0.8233 - val_loss: 0.4328

Epoch 88/100
89/89  0s 1ms/step - accuracy: 0.8115 - loss: 0.4305 - val_accuracy: 0.8233 - val_loss: 0.4331

Epoch 89/100
89/89  0s 1ms/step - accuracy: 0.8137 - loss: 0.4371 - val_accuracy: 0.8219 - val_loss: 0.43

```

32
Epoch 90/100
89/89 ————— 0s 1ms/step - accuracy: 0.8034 - loss: 0.4469 - val_accuracy: 0.8190 - val_loss: 0.43
32
Epoch 91/100
89/89 ————— 0s 1ms/step - accuracy: 0.8225 - loss: 0.4202 - val_accuracy: 0.8197 - val_loss: 0.43
42
Epoch 92/100
89/89 ————— 0s 1ms/step - accuracy: 0.8103 - loss: 0.4371 - val_accuracy: 0.8204 - val_loss: 0.43
36
Epoch 93/100
89/89 ————— 0s 1ms/step - accuracy: 0.8165 - loss: 0.4259 - val_accuracy: 0.8211 - val_loss: 0.43
36
Epoch 94/100
89/89 ————— 0s 1ms/step - accuracy: 0.8158 - loss: 0.4313 - val_accuracy: 0.8190 - val_loss: 0.43
37
Epoch 95/100
89/89 ————— 0s 1ms/step - accuracy: 0.8204 - loss: 0.4150 - val_accuracy: 0.8190 - val_loss: 0.43
35
Epoch 96/100
89/89 ————— 0s 1ms/step - accuracy: 0.8189 - loss: 0.4381 - val_accuracy: 0.8183 - val_loss: 0.43
40
Epoch 97/100
89/89 ————— 0s 1ms/step - accuracy: 0.8028 - loss: 0.4439 - val_accuracy: 0.8183 - val_loss: 0.43
35
Epoch 98/100
89/89 ————— 0s 1ms/step - accuracy: 0.8246 - loss: 0.4132 - val_accuracy: 0.8204 - val_loss: 0.43
37
Epoch 99/100
89/89 ————— 0s 1ms/step - accuracy: 0.8289 - loss: 0.4131 - val_accuracy: 0.8176 - val_loss: 0.43
46
Epoch 100/100
89/89 ————— 0s 1ms/step - accuracy: 0.8200 - loss: 0.4261 - val_accuracy: 0.8169 - val_loss: 0.43
41

```

```

In [46]: # Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc:.4f}")

```


```

45/45 ————— 0s 417us/step - accuracy: 0.8044 - loss: 0.4370
Test Accuracy: 0.8169


```

```
In [47]: # Plot training & validation accuracy values
plt.figure(figsize=(12, 7))
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.figure(figsize=(12, 7))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```




Model loss



```
In [48]: y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5).astype(int)

print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

45/45  0s 986us/step

```
[[ 934  102]
 [ 156  217]]
```

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1036
1	0.68	0.58	0.63	373
accuracy			0.82	1409
macro avg	0.77	0.74	0.75	1409
weighted avg	0.81	0.82	0.81	1409

Key Findings

Summary of Findings

1.High Monthly Charges: Customers with higher monthly charges are more likely to churn. This suggests a need for value

reassessment or targeted discounts for high-paying customers.

2.Contract Type: Customers on short-term or month-to-month contracts are more likely to churn. Incentivizing longer-term contracts could improve retention.

3.Service Usage: Customers who do not use additional services (e.g., MultipleLines, OnlineSecurity) may be more prone to churn. Offering bundled services at a discount might help retain these customers.

4.Proactive Interventions: By predicting churn, the company can proactively reach out to at-risk customers before they decide to leave, potentially reducing overall churn rates.

5.Tenure and Loyalty:Customers with shorter tenure (e.g., less than a year) are often at higher risk of churn. This suggests that early engagement strategies are crucial. Companies could implement loyalty programs, onboarding processes, or personalized communication to foster stronger relationships with new customers and reduce the likelihood of early churn.

6.Payment Method:The model might reveal that customers using certain payment methods (e.g., month-to-month billing with manual payments) have a higher churn rate. Encouraging customers to switch to automated payments or offering discounts for upfront payments could reduce churn by increasing convenience and commitment.

7.Customer Support Interaction:Frequent interactions with customer support, especially those involving complaints, might indicate dissatisfaction, leading to a higher risk of churn. Improving customer support, resolving issues promptly, and offering compensation for negative experiences could enhance customer satisfaction and retention.

8.Internet Service Type:If certain internet service types (e.g., DSL) are associated with higher churn rates compared to others (e.g., fiber optic), this could indicate a need to upgrade infrastructure or offer better service plans to customers in areas with inferior service quality.

9.Promotion and Discount Utilization:Customers who initially signed up during promotional periods or with discounts might have higher churn rates once those promotions expire. To retain these customers, the company could offer extended promotions, loyalty discounts, or alternative value-added services when the initial offer ends.

10.Geographical Location:If the model indicates that churn rates are higher in specific geographical areas, it may suggest local competition, service quality issues, or demographic factors. Targeted marketing campaigns, infrastructure improvements, or

localized customer engagement efforts could help address these issues.

11. Service Upgrades and Downgrades: Customers who frequently change their service plans (e.g., downgrading to a lower tier) may be at higher risk of churn. This could suggest dissatisfaction with the value received. Offering more flexible plans or value-added features that encourage customers to stay could mitigate this risk.

Strategic Recommendations

1. Early Engagement Programs: Implement initiatives such as personalized welcome packages, early-stage loyalty rewards, or tailored communication strategies to improve retention among new customers.
2. Automated Payment Incentives: Offer discounts or additional benefits to customers who switch to automated payment methods to reduce the risk of churn due to payment convenience.
3. Customer Feedback Loops: Establish regular feedback loops, particularly for customers who frequently contact support, to address pain points and enhance satisfaction.
4. Service Upgrade Incentives: Create targeted campaigns that encourage customers to upgrade to higher-tier services, offering additional features or benefits to improve perceived value.
5. Localized Marketing: Develop geographically targeted marketing efforts to address specific regional churn trends, such as improved infrastructure or localized service offerings.